


Linear classification models: Perceptron

Basic concepts:

- The Perceptron algorithm
- Perceptron loss/ hinge loss
- Subgradient descent
- Convergence proof of Perceptron
- Concept of Margin
- Voted and average Perceptrons

Linear Classifier



- We have discussed Logistic Regression
- LR learns $p(y|\mathbf{x})$:

$$P(y = 1|\mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{w}^T \mathbf{x})}$$

which yields a linear decision boundary

$$\mathbf{w}^T \mathbf{x} = 0$$

- We will now look at a different paradigm for learning a linear decision boundary directly
 - Without a probabilistic model
 - Based on a loss function

Binary classification: General Setup

- Given a set of training examples $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$, where each $\mathbf{x}_i \in R^d$, $y_i \in \{-1, 1\}$

- Learn a linear function

$$g(\mathbf{x}, \mathbf{w}) = w_0 + w_1 x_1 + \dots + w_d x_d$$

Given an example $\mathbf{x} = [x_1, \dots, x_d]^T$:

- predict $y(\mathbf{x}) = 1$ if $g(\mathbf{x}, \mathbf{w}) \geq 0$
 - predict $y(\mathbf{x}) = -1$ otherwise
- Compactly the classifier can be represented as:
 - $y(\mathbf{x}) = \text{sign}(w_0 + w_1 x_1 + \dots + w_d x_d) = \text{sign}(\mathbf{w}^T \mathbf{x})$
where $\mathbf{w} = [w_0, w_1, \dots, w_d]^T$, and $\mathbf{x} = [1, x_1, \dots, x_d]^T$
- Goal: find a good \mathbf{w} that minimizes some loss function $J(\mathbf{w})$

Loss function

$$J(\mathbf{w}) = \frac{1}{n} \sum_{m=1}^n L(g(\mathbf{w}, \mathbf{x}_m), y_m)$$


Where $L(g(\mathbf{w}, \mathbf{x}), y)$ is the loss of $g(\mathbf{w}, \mathbf{x})$ given its true label is y

0/1-loss:

$$L_{0/1}(g(\mathbf{w}, \mathbf{x}), y) = \begin{cases} 1 & \text{if } g(\mathbf{w}, \mathbf{x}) \text{ predicts wrong} \\ 0 & \text{otherwise} \end{cases}$$

This loss is conceptually aligned with our goal of maximizing accuracy, and minimizing error.

0/1 Loss counts the # of mistakes



Issue:

- Non convex – in general can be NP-hard to optimize
- Non-smooth – does not produce useful gradient since the surface of $J_{0/1}$ is **piece-wise flat**

Perceptron Loss


$$L_p(g(\mathbf{w}, \mathbf{x}), y) = \max(0, -y\mathbf{w}^T \mathbf{x})$$

- If prediction is correct, $-y\mathbf{w}^T \mathbf{x} < 0$,
$$L_p = \max(0, -y\mathbf{w}^T \mathbf{x}) = 0$$
- If prediction is incorrect, $-y\mathbf{w}^T \mathbf{x} > 0$,
$$L_p = \max(0, -y\mathbf{w}^T \mathbf{x}) = -y\mathbf{w}^T \mathbf{x} > 0$$
- Loss is a linear function of the weights


Perceptron Loss

$$J_p(w) = \frac{1}{n} \sum_{m=1}^n \max(0, -y_m \mathbf{w}^T \mathbf{x}_m)$$

- J_p is still non-smooth but piecewise linear
- Imagine if we drop a ball on this surface, it will with subgradient has a nice gradient leading to the solution region



0/1 loss: piecewise constant



Perceptron criterion: piecewise linear

Stochastic Gradient Descent

- The objective function consists of a sum over data points---
Stochastic Gradient Descent updates the parameter after observing each example


$$J(\mathbf{w}) = \frac{1}{n} \sum_{m=1}^n \max(0, -y_m \mathbf{w}^T \mathbf{x}_m)$$

$$J_m(\mathbf{w}) = \max(0, -y_m \mathbf{w}^T \mathbf{x}_m)$$

Subgradient:

$$\nabla J_m = \begin{cases} 0 & \text{if } y_m \mathbf{w} \cdot \mathbf{x}_m > 0 \\ -y_m \mathbf{x}_m & \text{otherwise} \end{cases}$$

$$\text{hinge}(t) = \max(0, h - t)$$



Perceptron
Hinge loss
with hinge=0


SVM
Hinge loss with
hinge=1

Perceptron Update Rule

After observing (\mathbf{x}_m, y_m) , if it is a mistake $\mathbf{w} \leftarrow \mathbf{w} + y_m \mathbf{x}_m$

Subgradient (subderivative)

$$\text{hinge}(t) = \max(0, -t)$$



Perceptron
Hinge loss
with hinge=0

Generalization of gradient to non-differentiable functions:

- Gradient of a function f can be viewed as defining a tangent line that touches f and lies below f
- Subgradient can be viewed as the set of lines that touches f and lies below f
 - At a differentiable point of the function it will be unique and the same as gradient
 - At non-smooth point, this will be a real set just pick one to work with
- Subgradient descent behaves similarly to gradient descent in terms of convergence

Stochastic Gradient Descent

- The objective function consists of a sum over data points---
Stochastic Gradient Descent updates the parameter after observing each example


$$J(\mathbf{w}) = \frac{1}{n} \sum_{m=1}^n \max(0, -y_m \mathbf{w}^T \mathbf{x}_m)$$

$$J_m(\mathbf{w}) = \max(0, -y_m \mathbf{w}^T \mathbf{x}_m)$$

$\text{hinge}(t)$ is not smooth, we compute its subgradient

$$\nabla J_m = \begin{cases} 0 & \text{if } y_m \mathbf{w} \cdot \mathbf{x}_m > 0 \\ -y_m \mathbf{x}_m & \text{otherwise} \end{cases}$$

$$\text{hinge}(t) = \max(0, h - t)$$



Perceptron
Hinge loss
with hinge=0

SVM
Hinge loss with
hinge=1

Perceptron Update Rule

After observing (\mathbf{x}_m, y_m) , if it is a mistake $\mathbf{w} \leftarrow \mathbf{w} + y_m \mathbf{x}_m$

The Perceptron Algorithm (online)

(Stochastic gradient descent)

Let $\mathbf{w} \leftarrow (0,0,0,\dots,0)$

Repeat until convergence

for every training example $m = 1, \dots, n$:

$$u_m \leftarrow \mathbf{w}^T \mathbf{x}_m$$

$$\text{if } y_m \cdot u_m \leq 0 \quad \mathbf{w} \leftarrow \mathbf{w} + y_m \mathbf{x}_m$$

Online

- Look at one example at a time, update the model as soon as we make an error – as opposed to batch algorithms that update parameters after seeing the entire training set.

Error-driven

- We only update parameters/model if we make an error

Effect of the perceptron update

- Our current weight is w_t and it makes a mistake on (x_i, y_i) , i.e., $y_i w_t^T x_i \leq 0$
- Perform update







$$w_{t+1} = w_t + y_i x_i$$

- What can we say about w_{t+1} 's performance on (x_i, y_i) ?

See concept warehouse

When an error is made, moves the weight in a direction that corrects the error

Update 1



Update 2

Convergence Theorem

(Block, 1962, Novikoff, 1962)

Given training example sequence $(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)$.

If $\forall i, \|\mathbf{x}_i\| \leq D$, and $\exists \mathbf{u}, \|\mathbf{u}\| = 1$ and $y_i \mathbf{u}^T \mathbf{x}_i \geq \gamma > 0$ for all i ,

then the number of mistakes that the perceptron algorithm makes is at most $(D / \gamma)^2$.

Note that $\|\cdot\|$ denotes the Euclidean norm of a vector.

Proof

Let \mathbf{u} be a unit vector that achieves γ margin, i.e., $|\mathbf{u}| = 1$ and $\forall i, y_i \mathbf{u}^T \mathbf{x}_i \geq \gamma$

Let \mathbf{x}_k be the k th mistake, we have $\mathbf{w}(k) = \mathbf{w}(k-1) + y_k \mathbf{x}_k$

Main idea: we want to show that the direction of \mathbf{w}_k converges to \mathbf{u} ,

i.e., $\frac{\mathbf{u}^T \mathbf{w}_k}{|\mathbf{u}| |\mathbf{w}_k|}$ converges to 1. To show this, we work out two parts:

1. $\mathbf{u}^T \mathbf{w}_k$ grows bigger as k increases
2. $|\mathbf{w}_k|$ does not grow as fast

Part 1:

$$\mathbf{u}^T \mathbf{w}_k = \mathbf{u}^T (\mathbf{w}_{k-1} + y_k \mathbf{x}_k) = \mathbf{u}^T \mathbf{w}_{k-1} + y_k \mathbf{u}^T \mathbf{x}_k \geq \mathbf{u}^T \mathbf{w}_{k-1} + \gamma \geq k\gamma$$

Part 2:

$$\begin{aligned} \mathbf{w}_k^T \mathbf{w}_k &= (\mathbf{w}_{k-1} + y_k \mathbf{x}_k)^T (\mathbf{w}_{k-1} + y_k \mathbf{x}_k) \\ &= \mathbf{w}_{k-1}^T \mathbf{w}_{k-1} + 2y_k \mathbf{w}_{k-1}^T \mathbf{x}_k + \mathbf{x}_k^T \mathbf{x}_k \leq \mathbf{w}_{k-1}^T \mathbf{w}_{k-1} + D^2 \leq kD^2 \end{aligned}$$

Putting it together: $\frac{\mathbf{u}^T \mathbf{w}_k}{|\mathbf{u}| |\mathbf{w}_k|} \geq \frac{k\gamma}{\sqrt{k}D}$, but this cannot exceed 1. so we must

have $\frac{k\gamma}{\sqrt{k}D} \leq 1 \Rightarrow$

$$k \leq \left(\frac{D}{\gamma} \right)^2$$

Margin

- γ is referred to as the **margin**
 - Min distance from data points to the decision boundary
 - Bigger margin -> easier classification problems
- This concept will be utilized later by support vector machines

Implications of the convergence theorem

- See concept warehouse

Practical considerations for online perceptron

- The order of training examples matters!
 - Random is better
- Given large amounts of training data, early stopping can be used to avoid overfitting
- Simple modification can significantly improve performance
 - Voted perceptron and average perceptron

Voted Perceptron

- Keep intermediate hypotheses and have them vote [Freund and Schapire 1998]

Let $\mathbf{w} \leftarrow (0,0,0,\dots,0)$

$c_0 = 0, n = 0$

Repeat for T times

for each training example i :

$$u_i \leftarrow \mathbf{w}^T \mathbf{x}_i$$

if $y_i u_i \leq 0$

$$\mathbf{w}_{n+1} \leftarrow \mathbf{w}_n + y_i \mathbf{x}_i$$

$$n = n + 1$$

$$c_n = 0$$

else $c_n = c_n + 1$

The output will be a collection of linear separators $\mathbf{w}_0, \mathbf{w}_1, \dots, \mathbf{w}_M$ along with their survival time c_0, c_1, \dots, c_M

The c 's can be viewed as measures of the reliability of the \mathbf{w} 's

For classification, take a weighted vote among all separators:

$$\hat{y} = \text{sign} \left\{ \sum_{n=0}^N c_n \text{sign}(\mathbf{w}_n^T \mathbf{x}) \right\}$$

Average Perceptron

- Voted perceptron requires storing all intermittent weights
 - Large memory consumption
 - Slow prediction time
- Average perceptron

$$\hat{y} = \text{sign}\left\{\left(\sum_{n=0}^N c_n \mathbf{w}_n^T\right) \mathbf{x}\right\}$$





- Take the weighted average of all the intermittent weights
- Can be implemented by maintaining a running average, no need to store all weights
- Fast prediction time

Final Discussion on Perceptron

- Learns $\hat{y} = f(\mathbf{x})$ directly – a **discriminative** method
- Performs stochastic (sub-)gradient descent to optimize the perceptron loss (also called hinge loss with hinge=0, will see more hinge loss later in SVM)
- Guaranteed to converge in finite steps if the data is linearly separable
 - # of updates is inversely proportional to the **margin** of the optimal decision boundary
 - Guarantees convergence but not necessarily to the maximum margin separator – again we will address this later in SVM
- Voted and average perceptrons provide significant performance improvement in practice

Beyond the Basic Perceptron

Structured Prediction with Perceptrons



...

A general problem

- Given some **input x**
 - An email, a sentence ...
- Consider a set of **candidate outputs y**
 - **Classifications** for x (small number: often just 2)
 - **Taggings** of x (exponentially many)
 - **Parses** of x (exponentially many)
 - **Translations** of x (exponentially many)
 - ...
- Want to **find the “best” y , given x**

Structured prediction

Scoring by Linear Models

- Given some **input** x
- Consider a set of **candidate outputs** y
- Define a scoring function $\text{score}(x, y)$

Linear function: A sum of feature weights (you pick the features!)

Weight of feature k
(learned or set by hand)

$$\text{score}(x, y) = \sum_k \theta_k f_k(x, y)$$

Ranges over all features,
e.g., $k=5$ (numbered features)
or $k=\text{"see Det Noun"}$ (named features)

Whether (x, y) has feature k (0 or 1)
Or how many times it fires (≥ 0)
Or how strongly it fires (real #)

- Choose y that maximizes $\text{score}(x, y)$

Scoring by Linear Models

- Given some **input x**
- Consider a set of **candidate outputs y**
- Define a scoring function $\text{score}(x, y)$

Linear function: A sum of feature weights (you pick the features!)

(learned or set by hand)

$$\text{score}(x, y) = \boxed{\vec{\theta}} \cdot \vec{f}(x, y)$$

This linear decision rule is sometimes called a “perceptron.”
It’s a “structured perceptron” if it does structured prediction
(number of y candidates is unbounded, e.g., grows with $|x|$).

- Choose **y** that maximizes $\text{score}(x, y)$

Perceptron Training Algorithm

- initialize θ (usually to the zero vector)
- repeat:
 - Pick a training example (x, y)
 - Model predicts y^* that maximizes $\text{score}(x, y^*)$
 - Update weights by a step of size $\varepsilon > 0$:
$$\theta = \theta + \varepsilon \cdot (f(x, y) - f(x, y^*))$$

If model prediction was wrong ($y \neq y^*$), then we must have $\text{score}(x, y) \leq \text{score}(x, y^*)$ instead of $>$ as we want.

Equivalently, $\theta \cdot f(x, y) \leq \theta \cdot f(x, y^*)$

Equivalently, $\theta \cdot (f(x, y) - f(x, y^*)) \leq 0$ but we want it positive.

Our update increases it (by $\varepsilon \cdot ||f(x, y) - f(x, y^*)||^2 \geq 0$)

Perceptron for Structured Prediction

- What we see here is the same as the regular perceptron
- Similar convergence guarantee
- The challenge is the inference part
 - Finding the y that maximizes the score for given x
 - Cannot resort to brute-force enumeration
 - Much research goes into
 - How to devise proper features and efficient algorithms for inference
 - How to perform approximate inference
 - How to learn when inference is approximate